

Lab (2)

Name	Sec	BN
Basma Hatem Elhoseny	1	16
Sarah Mohamed Hossam Hassan	1	29

Requirement(1) [Matrix Addition]:

Kernel (1) Each Thread Produces one output matrix element:

$$GridSize = (ceil(\frac{No.Cols - 1}{16}), ceil(\frac{No.rows - 1}{16}))$$

Case(1) 3*4 Matrix

With configuration that kernel size 16*16 and one thread per element so:

```
3-4
Total no of blocks 1
Total no of threads 256
```

```
==18272== Profiling application: out_1.exe ./tests/test_3_4.txt out.txt
==18272== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 93.78% 40.516us      1 40.516us 40.516us 40.516us add_matrix(float*, float*, float*, int, int)
                3.85% 1.6640us      2      832ns 352ns 1.3120us [CUDA memcpy HtoD]
                2.37% 1.0240us      1 1.0240us 1.0240us 1.0240us [CUDA memcpy DtoH]
API calls: 90.78% 473.60ms      3 157.87ms 3.5000us 473.59ms cudaMalloc
                5.72% 29.824ms      1 29.824ms 29.824ms 29.824ms cuDevicePrimaryCtxRelease
                3.19% 16.653ms      3 5.5509ms 48.400us 16.337ms cudaMemcpy
                0.12% 627.40us      3 209.13us 5.6000us 559.70us cudaFree
                0.08% 437.10us      1 437.10us 437.10us 437.10us cuLibraryUnload
                0.08% 415.80us      1 415.80us 415.80us 415.80us cuLibraryLoadData
                0.02% 99.600us      1 99.600us 99.600us 99.600us cudaLaunchKernel
                0.00% 25.700us     114      225ns 100ns 2.2000us cuDeviceGetAttribute
                0.00% 9.2000us      3 3.0660us 300ns 8.4000us cuDeviceGetCount
                0.00% 4.5000us      2 2.2500us 200ns 4.3000us cuDeviceGet
                0.00% 3.2000us      1 3.2000us 3.2000us 3.2000us cuModuleGetLoadingMode
                0.00% 800ns      1      800ns 800ns 800ns cuDeviceGetName
                0.00% 300ns      1      300ns 300ns 300ns cuDeviceTotalMem
                0.00% 300ns      1      300ns 300ns 300ns cuDeviceGetLuid
                0.00% 200ns      1      200ns 200ns 200ns cuDeviceGetUuid
```

D:\Parallel Computing Labs\Lab2\requirement>

Kernel(2) Each Thread Produces one output matrix row:

$$GridSize = \text{ceil}\left(\frac{No.Rows - 1}{16}\right)$$

Case(1) 3*4 Matrix

With configuration that kernel size 256 and one thread per rows so:

```
rows:3-cols:4
Total no of blocks 1
Total no of threads 256
```

```
==21016== Profiling application: out_2.exe ./tests/test_3_4.txt out.txt
==21016== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 98.47% 167.09us      1 167.09us 167.09us 167.09us add_matrix(float*, float*, float*, int, int)
                0.98% 1.6640us      2    832ns    352ns 1.3120us [CUDA memcpy HtoD]
                0.55%  928ns      1    928ns    928ns  928ns [CUDA memcpy DtoH]
API calls: 67.85% 142.81ms      3 47.604ms 3.8000us 142.80ms cudaMalloc
            18.65% 39.247ms      1 39.247ms 39.247ms 39.247ms cuDevicePrimaryCtxRelease
            12.70% 26.727ms      3 8.9090ms 41.100us 26.400ms cudaMemcpy
            0.28% 586.70us      3 195.57us 6.1000us 533.70us cudaFree
            0.21% 438.20us      1 438.20us 438.20us 438.20us cuLibraryLoadData
            0.20% 415.20us      1 415.20us 415.20us 415.20us cuLibraryUnload
            0.10% 216.40us      1 216.40us 216.40us 216.40us cudaLaunchKernel
            0.01% 26.800us     114    235ns    100ns 2.2000us cuDeviceGetAttribute
            0.00% 6.6000us      2 3.3000us    300ns 6.3000us cuDeviceGet
            0.00% 5.9000us      3 1.9660us    300ns 5.0000us cuDeviceGetCount
            0.00% 3.0000us      1 3.0000us 3.0000us 3.0000us cuModuleGetLoadingMode
            0.00%  800ns      1    800ns    800ns  800ns cuDeviceGetName
            0.00%  400ns      1    400ns    400ns  400ns cuDeviceTotalMem
            0.00%  400ns      1    400ns    400ns  400ns cuDeviceGetLuid
            0.00%  200ns      1    200ns    200ns  200ns cuDeviceGetUuid

D:\Parallel Computing Labs\Lab2\requirement>
```

Kernel(3) Each Thread Produces one output matrix column:

$$GridSize = \text{ceil}\left(\frac{No.Rows - 1}{16}\right)$$

Case(1) 3*4 Matrix

With configuration that kernel size 256 and one thread per rows so:

```
Total no of blocks 1
Total no of threads 256
```

```

==4172== Profiling application: out_3.exe ./tests/test_3_4.txt out.txt
==4172== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 97.54% 102.96us    1 102.96us 102.96us 102.96us add_matrix(float*, float*, float*, int, int)
               1.55% 1.6320us    2   816ns   320ns  1.3120us [CUDA memcpy HtoD]
               0.91% 961ns      1   961ns   961ns   961ns [CUDA memcpy DtoH]
API calls: 68.63% 135.98ms    3 45.327ms 5.9000us 135.94ms cudaMalloc
            18.66% 36.968ms    1 36.968ms 36.968ms 36.968ms cuDevicePrimaryCtxRelease
            11.90% 23.590ms    3 7.8632ms 59.200us 23.255ms cudaMemcpy
            0.29% 571.40us    3 190.47us 6.2000us 534.40us cudaFree
            0.22% 440.20us    1 440.20us 440.20us 440.20us cuLibraryLoadData
            0.16% 325.50us    1 325.50us 325.50us 325.50us cuLibraryUnload
            0.12% 229.40us    1 229.40us 229.40us 229.40us cudaLaunchKernel
            0.01% 28.500us   114 250ns    100ns  2.2000us cuDeviceGetAttribute
            0.00% 6.1000us    3 2.0330us 300ns   5.1000us cuDeviceGetCount
            0.00% 5.7000us    2 2.8500us 300ns   5.4000us cuDeviceGet
            0.00% 3.3000us    1 3.3000us 3.3000us 3.3000us cuModuleGetLoadingMode
            0.00% 800ns      1 800ns    800ns   800ns cuDeviceGetName
            0.00% 400ns      1 400ns    400ns   400ns cuDeviceTotalMem
            0.00% 300ns      1 300ns    300ns   300ns cuDeviceGetLuid
            0.00% 200ns      1 200ns    200ns   200ns cuDeviceGetUuid

D:\Parallel Computing Labs\Lab2\requirement>_

```

Comments:

- ✓ Case 3*4 it is clear that the kernel 1 is the fastest regarding computing the addition function and nearly the 3 kernels have near copying time form host to device and vice versa 😊
- ✓ We think whatever the matrix size, kernel (1) will be the fastest Are we right ?! 😊 Let's see <3

Bench Marking:

Notice that these results are based on running matrix produced as random numbers generated by a script :D. [Runed on Colab]

Matrix Shape	Kernel(1) [element]	Kernel(2) [row]	Kernel(3) [col]
3x4	80.672us	279.52us	213.41us
2x2	80.159us	143.94us	144.10us
2x4	80.480us	277.12us	145.79us
Total of 10,000 elements			
100x100	33.184ms	14.142ms	13.420ms
50x200	31.617ms	17.931ms	10.002ms
200x50	28.655ms	10.427ms	17.857ms
10000x1	32.122ms	30.463ms	448.05ms
1x10000	24.259ms	286.08ms	33.404ms
Total of 50,000 elements			
50000x1	238.71ms	1.16151s	1.56368s
1x50000	318.95ms	1.43044s	1.03092s
1000x50	1.54740s	44.100ms	89.304ms
50x1000	1.10863s	87.770ms	43.907ms
Total of 1,000,000 elements			

1000x1000	133.538s	882.08ms	876.51ms
1000000x1	4.85198s	114.753s	28.5801s
1x1000000	6.47006s	28.4147s	141.907s

Notes:

For the First Three example since matrix dim is very small we don't sense a lot of improvement between the 3 kernels so we will apply like stress test 🇧🇪🇧🇪

Total of 10,000-element Matrix

For 100x100:

- ✓ The three kernels are nearly the same except that kernel(1) is bit higher due to more threads are required 100x100 while in both kernel(2)&(3) only 100 thread is required.

For 50x200:

- ✓ Kernel(3) is the best bec only 200 thread is required each thread make only 50 operation :D

For 200x50:

- ✓ Kernel(2) is the best bec only 200 thread is required each thread make only 50 operation :D

For 10000x1:

- ✓ Worst Is Kernel(3) 448.05 ms because simply this is computed by single thread (No parallelization) :D

For 1x10000:

- ✓ Worst Is Kernel(2) 286.08 ms because simply this is computed by single thread (No parallelization) :D

Total of 1,000,000-element Matrix

For 1000x1000:

- ✓ It is clear that the worst Time is for Kernel(1) where each thread is computing 1 element but since here the large no of elements we need 1000x1000 thread in total which may not be available to be used together so some is computed with the available threads and finishes then others (Scarcity of Thread compared to the huge no of threads required)
- ✓ While we get better time for kernel(2) & kernel(3) due to less no of threads required in both cases [1000 only :D] (This proves that our claim above is Wrong: kernel(1) is always the best ~~XXX~~)

For 1x1000000 & 1000000x1:

- ✓ The Best will be the one will the orientation corresponding to the no of threads and the other will be the worst because on thread is responsible for the computation (No parallelism)

Conclusion:

We have corrected our faulty claim that kernel(1) will be always the same :D. It depends according to the total no of elements.

But it is clear that of course if the no of rows are more then kernel(2) is the best and kernel(3) will be the worst. And the same for the cols case .

That's All 😁